

# **INFO-6147 Deep Learning With Pytorch**

## **Fall 2024**

### **Project Final Report**

---

**Proposal: Tiny ImageNet Classification**

**Name: Mohamed Hossni Ali**

**ID: 1306850**

## Contents

Table of Figures .....	2
Introduction.....	3
Objective .....	3
Dataset Description .....	3
Data Preprocessing.....	4
Methodology.....	4
Classifier 1 Architecture .....	4
Classifier 2 Architecture .....	5
Classifier 3 Architecture .....	5
Classifier 4 Architecture .....	6
Classifier 4.2 Architecture.....	7
Classifier 5 Architecture .....	8
Classifier 7 Architecture .....	9
Data Augmentation Methods: .....	9
Switching off pixels .....	10
Creating new images.....	10
Results .....	11
Conclusions and recommendations .....	14
Project Hosting: .....	14
Final Notes .....	14
References: .....	15
Appendix A: Tiny ImageNet Classes .....	16

## Table of Figures

Figure 1: Tiny ImageNet sample .....	3
Figure 2: Classifier 1 architecture.....	4
Figure 3: Classifier 2 architecture.....	5
Figure 4: Classifier 3 architecture.....	6
Figure 5: Classifier 4 architecture.....	7
Figure 6: Classifier 4.2 architecture with return x+y.....	7
Figure 7: Classifier 5 architecture.....	8
Figure 8: Classifier 7 architecture.....	9
Figure 9: Data Augmentation, random switching off of pixel (noise) .....	10
Figure 10: Data augmentation, rotation and flipping.....	11
Figure 11: Training loss chart for all the models.....	12
Figure 12: test loss for all models.....	12
Figure 13: Models accuracy over epochs .....	13
Figure 14: resnet18 pretrained model accuracy .....	13
Figure 15: accuracy comparison for all the models .....	14

## Introduction

Image classification tasks are important part of computer vision which is used in various applications. This project aims to create an image classification model using convolutional neural network (CNN) and then compare its results with state of the art models available on the web. The dataset that is used in this project is Tiny ImageNet.

## Objective

The objective of this project is to create different classifiers with reasonable accuracy on Tiny ImageNet dataset. Additionally, transfer learning will be used to compare the model which is made from scratch to the industry state of the art models.

## Dataset Description

Tiny ImageNet dataset is a group of colored images with 200 classes. The dataset contains 100,000 samples with balanced distribution of 500 images per class for training (50 for validation and 50 for testing). The images in this dataset are of size 64X64X3. The list of classes are presented in Appendix A [1].

The data comes in zip file which when extracted explodes multiple folders. One folder for training contains subfolders for each class which has images folder for images. The validation set is contained in one val folder that has the 10k validation images and text file that contain the class reference of these images.

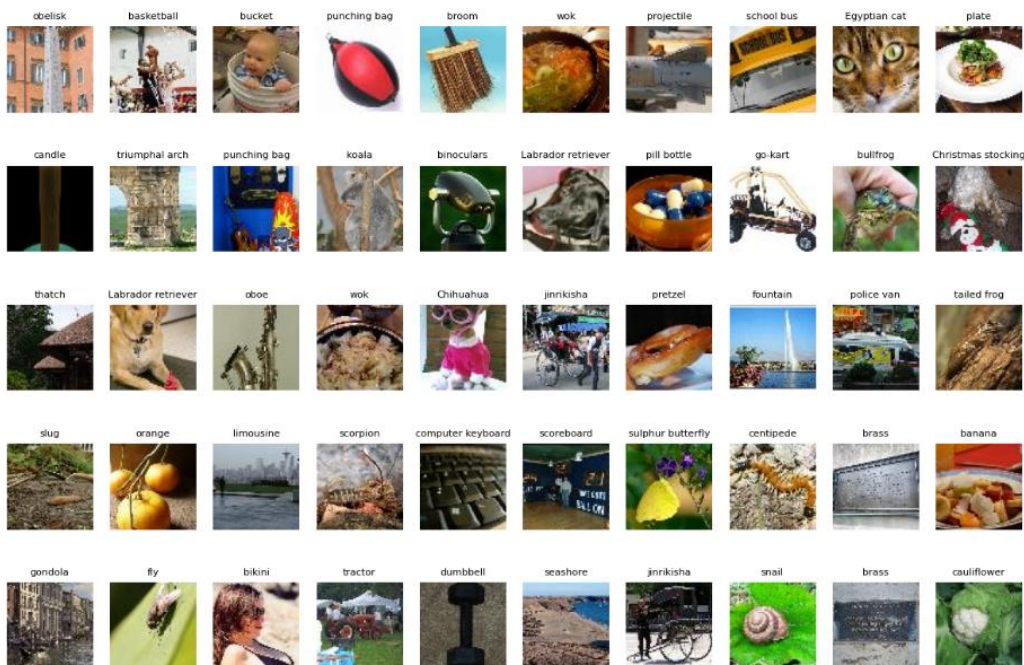


Figure 1: Tiny ImageNet sample

## Data Preprocessing

Before training the model with the dataset the mean and standard deviation (std) values were calculated on the trainset for normalization. The mean was [0.4802, 0.4481, 0.3975] and the standard deviation [0.2296, 0.2263, 0.2255] for the three channels.

The ImageFolder function was used to create the training set data loader while a custom function was used to create the validation dataset loader as the labels had to be linked differently.

The final transformation included normalization and conversion to tensor.

## Methodology

The methodology for classification in this project is conducted using CNN. CNNs can capture spatial information with filters and they would be the appropriate candidates for this task.

Multiple network designs were constructed and trained with the dataset. The following section describes each of these designs in details.

### Classifier 1 Architecture

The first model constitutes of 8 convolutional layers and 3 fully connected layers at the end. ReLu activation function is applied after every convolution and max pooling is done after every two consecutive convolutions to reduce the size. A regulation was applied throughout by using dropout layers.

```
class Classifier(nn.Module):
    def __init__(self):
        super(Classifier, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, padding=1)
        self.BN1 = nn.BatchNorm2d(num_features=32)
        self.conv2 = nn.Conv2d(in_channels=32, out_channels=32, kernel_size=3, padding=1)
        self.BN2 = nn.BatchNorm2d(num_features=32)

        self.conv2_drop = nn.Dropout2d()

        self.conv3 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, padding=1)
        self.BN3 = nn.BatchNorm2d(num_features=64)
        self.conv4 = nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3, padding=1)
        self.BN4 = nn.BatchNorm2d(num_features=64)

        self.conv5 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, padding=1)
        self.BN5 = nn.BatchNorm2d(num_features=128)
        self.conv6 = nn.Conv2d(in_channels=128, out_channels=128, kernel_size=3, padding=1)
        self.BN6 = nn.BatchNorm2d(num_features=128)

        self.conv7 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3, padding=1)
        self.BN7 = nn.BatchNorm2d(num_features=256)
        self.conv8 = nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3, padding=1)
        self.BN8 = nn.BatchNorm2d(num_features=256)

        self.fc1 = nn.Linear(4096, 1024)
        self.fc2 = nn.Linear(1024, 512)
        self.fc3 = nn.Linear(512, 200)

    def forward(self, x):
        x = F.relu(self.BN1(self.conv1(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.conv2(x)), 2))
        x = self.conv2_drop(F.relu(self.BN3(self.conv3(x))))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN4(self.conv4(x))), 2))
        x = F.relu(self.BN5(self.conv5(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN6(self.conv6(x))), 2))
        x = F.relu(self.BN7(self.conv7(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN8(self.conv8(x))), 2))
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

Figure 2: Classifier 1 architecture

## Classifier 2 Architecture

The second classification model was similar to the first except there are middle fully connected layers that get reshaped back into the correct dimensions. The philosophy behind this was to change the output of the convolutions with weights.

```
class Classifier_middle_FC(nn.Module):
    def __init__(self):
        super(Classifier_middle_FC, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, padding=1)
        self.BN1 = nn.BatchNorm2d(num_features=32)
        self.conv2 = nn.Conv2d(in_channels=32, out_channels=32, kernel_size=3, padding=1)
        self.BN2 = nn.BatchNorm2d(num_features=32)
        self.conv2_drop = nn.Dropout2d()

        self.conv3 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, padding=1)
        self.BN3 = nn.BatchNorm2d(num_features=64)
        self.conv4 = nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3, padding=1)
        self.BN4 = nn.BatchNorm2d(num_features=64)

        self.fc_inner_2 = nn.Linear(16384, 16384)

        self.conv5 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, padding=1)
        self.BN5 = nn.BatchNorm2d(num_features=128)
        self.conv6 = nn.Conv2d(in_channels=128, out_channels=128, kernel_size=3, padding=1)
        self.BN6 = nn.BatchNorm2d(num_features=128)

        self.fc_inner_3 = nn.Linear(8192, 8192)

        self.conv7 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3, padding=1)
        self.BN7 = nn.BatchNorm2d(num_features=256)
        self.conv8 = nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3, padding=1)
        self.BN8 = nn.BatchNorm2d(num_features=256)

        self.fc1 = nn.Linear(4096, 1024)
        self.fc2 = nn.Linear(1024, 512)
        self.fc3 = nn.Linear(512, 200)

    def forward(self, x):
        x = F.relu(self.BN1(self.conv1(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.conv2(x)), 2))

        x = self.conv2_drop(F.relu(self.BN3(self.conv3(x))))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN4(self.conv4(x))), 2))
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc_inner_2(x))
        x = x.view(x.size(0), 64, 16, 16)

        x = F.relu(self.BN5(self.conv5(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN6(self.conv6(x))), 2))
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc_inner_3(x))
        x = x.view(x.size(0), 128, 8, 8)

        x = F.relu(self.BN7(self.conv7(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN8(self.conv8(x))), 2))
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

Figure 3: Classifier 2 architecture

## Classifier 3 Architecture

This model was designed with an intermediate linear layer that reduces the dimensions of the image before going into the last convolutional block. This was used as a replacement for max-pooling after the sixth convolutional block. The concept assumed that having weights that will reduce the dimension might work better than taking the max value using the max-pool method.

```

class Classifier_FC_asPool(nn.Module):
    def __init__(self):
        super(Classifier_FC_asPool, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, padding=1)
        self.BN1 = nn.BatchNorm2d(num_features=32)
        self.conv2 = nn.Conv2d(in_channels=32, out_channels=32, kernel_size=3, padding=1)
        self.BN2 = nn.BatchNorm2d(num_features=32)

        self.conv2_drop = nn.Dropout2d()

        self.conv3 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, padding=1)
        self.BN3 = nn.BatchNorm2d(num_features=64)
        self.conv4 = nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3, padding=1)
        self.BN4 = nn.BatchNorm2d(num_features=64)

        self.conv5 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, padding=1)
        self.BN5 = nn.BatchNorm2d(num_features=128)
        self.conv6 = nn.Conv2d(in_channels=128, out_channels=128, kernel_size=3, padding=1)
        self.BN6 = nn.BatchNorm2d(num_features=128)

        self.fc_inner_3 = nn.Linear(32768, 8192)

        #reshape here again
        self.conv7 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3, padding=1)
        self.BN7 = nn.BatchNorm2d(num_features=256)
        self.conv8 = nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3, padding=1)
        self.BN8 = nn.BatchNorm2d(num_features=256)

        self.fc1 = nn.Linear(4096, 1024)
        self.fc2 = nn.Linear(1024, 512)
        self.fc3 = nn.Linear(512, 200)

    def forward(self, x):
        x = F.relu(self.BN1(self.conv1(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN1(self.conv2(x)))), 2))

        x = self.conv2_drop(F.relu(self.BN3(self.conv3(x))))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN4(self.conv4(x)))), 2))

        x = F.relu(self.BN5(self.conv5(x)))
        x = self.conv2_drop(F.relu(self.BN6(self.conv6(x))))
        #reshaping to feed to FC_inner_3
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc_inner_3(x))
        x = x.view(x.size(0), 128, 8, 8)

        x = F.relu(self.BN7(self.conv7(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN8(self.conv8(x)))), 2))
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

```

Figure 4: Classifier 3 architecture

## Classifier 4 Architecture

This model has deeper convolutions (i.e each convolutional layer has more channels than the previous models). The first convolution starts with 128 channels and goes up to 1024 in contrast to the earlier model which had maximum of 256 channels in the last conv layer.

```
class Classifier_deep_conv5(nn.Module):
    def __init__(self):
        super(Classifier_deep_conv5, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3,out_channels=128,kernel_size=3, padding=1)
        self.BN1 = nn.BatchNorm2d(num_features=128)
        self.conv2 = nn.Conv2d(in_channels=128, out_channels=128, kernel_size=3, padding=1)
        self.BN2 = nn.BatchNorm2d(num_features=128)

        self.conv2_drop = nn.Dropout2d()

        self.conv3 = nn.Conv2d(in_channels=128,out_channels=256,kernel_size=3, padding=1)
        self.BN3 = nn.BatchNorm2d(num_features=256)
        self.conv4 = nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3, padding=1)
        self.BN4 = nn.BatchNorm2d(num_features=256)

        self.conv5 = nn.Conv2d(in_channels=256,out_channels=512,kernel_size=3, padding=1)
        self.BN5 = nn.BatchNorm2d(num_features=512)
        self.conv6 = nn.Conv2d(in_channels=512,out_channels=512,kernel_size=3, padding=1)
        self.BN6 = nn.BatchNorm2d(num_features=512)

        self.conv7 = nn.Conv2d(in_channels=512,out_channels=1024,kernel_size=3, padding=1)
        self.BN7 = nn.BatchNorm2d(num_features=1024)
        self.conv8 = nn.Conv2d(in_channels=1024,out_channels=1024,kernel_size=3, padding=1)
        self.BN8 = nn.BatchNorm2d(num_features=1024)

        self.fc1 = nn.Linear(16384,8192)
        self.fc2 = nn.Linear(8192,4096)
        self.fc3 = nn.Linear(4096, 2048)
        self.fc4 = nn.Linear(2048, 1024)
        self.fc5 = nn.Linear(1024,512)
```

Figure 5: Classifier 4 architecture

## Classifier 4.2 Architecture

Taking the previous model a step further and being inspired by resnet, this model has an inner linear layers that act as middle classifier. The value output from the middle classifier (y) is then summed with the final x value to boost the classification.

```
class Classifier_deep_conv5(nn.Module):
    def __init__(self):
        super(Classifier_deep_conv5, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3,out_channels=128,kernel_size=3, padding=1)
        self.BN1 = nn.BatchNorm2d(num_features=128)
        self.conv2 = nn.Conv2d(in_channels=128, out_channels=128, kernel_size=3, padding=1)
        self.BN2 = nn.BatchNorm2d(num_features=128)

        self.conv2_drop = nn.Dropout2d()

        self.conv3 = nn.Conv2d(in_channels=128,out_channels=256,kernel_size=3, padding=1)
        self.BN3 = nn.BatchNorm2d(num_features=256)
        self.conv4 = nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3, padding=1)
        self.BN4 = nn.BatchNorm2d(num_features=256)

        self.conv5 = nn.Conv2d(in_channels=256,out_channels=512,kernel_size=3, padding=1)
        self.BN5 = nn.BatchNorm2d(num_features=512)
        self.conv6 = nn.Conv2d(in_channels=512,out_channels=512,kernel_size=3, padding=1)
        self.BN6 = nn.BatchNorm2d(num_features=512)

        self.fc_mid1 = nn.Linear(32768,8192)
        self.fc_mid2 = nn.Linear(8192,4096)
        self.fc_mid3 = nn.Linear(4096, 2048)
        self.fc_mid4 = nn.Linear(2048, 1024)
        self.fc_mid5 = nn.Linear(1024,512)
        self.fc_mid6 = nn.Linear(512,200)

        self.conv7 = nn.Conv2d(in_channels=512,out_channels=1024,kernel_size=3, padding=1)
        self.BN7 = nn.BatchNorm2d(num_features=1024)
        self.conv8 = nn.Conv2d(in_channels=1024,out_channels=1024,kernel_size=3, padding=1)
        self.BN8 = nn.BatchNorm2d(num_features=1024)

        self.fc1 = nn.Linear(16384,8192)
        self.fc2 = nn.Linear(8192,4096)
        self.fc3 = nn.Linear(4096, 2048)
        self.fc4 = nn.Linear(2048, 1024)
        self.fc5 = nn.Linear(1024,512)
        self.fc6 = nn.Linear(512,200)

    def forward(self, x):
        x = F.relu(self.BN1(self.conv1(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.conv2(x)),2))

        x = self.conv2_drop(F.relu(self.BN3(self.conv3(x))))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN4(self.conv4(x))),2))

        x = F.relu(self.BN5(self.conv5(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN6(self.conv6(x))),2))
        y = x.view(x.size(0),-1)
        y = F.relu(self.fc_mid1(y))
        y = F.relu(self.fc_mid2(y))
        y = F.relu(self.fc_mid3(y))
        y = F.relu(self.fc_mid4(y))
        y = F.relu(self.fc_mid5(y))
        y = self.fc_mid6(y)

        x = F.relu(self.BN7(self.conv7(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.BN8(self.conv8(x))), 2))
        x = x.view(x.size(0),-1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = F.relu(self.fc3(x))
        x = F.relu(self.fc4(x))
        x = F.relu(self.fc5(x))
        x = self.fc6(x)

        return x+y
```

Figure 6: Classifier 4.2 architecture with return x+y



## Classifier 5 Architecture

This classifier adds more convolutions to the previous classifiers and has 16 convolutional layers with large number of channels like the classifier 4.

```
class Classifier_multi_convs(nn.Module):
    def __init__(self):
        super(Classifier_multi_convs, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, padding=1)
        self.BN1 = nn.BatchNorm2d(num_features=32)
        self.conv2 = nn.Conv2d(in_channels=32, out_channels=32, kernel_size=3, padding=1)
        self.BN2 = nn.BatchNorm2d(num_features=32)
        self.conv3 = nn.Conv2d(in_channels=32, out_channels=32, kernel_size=3, padding=1)
        self.BN3 = nn.BatchNorm2d(num_features=32)
        self.conv4 = nn.Conv2d(in_channels=32, out_channels=32, kernel_size=3, padding=1)
        self.BN4 = nn.BatchNorm2d(num_features=32)

        self.conv2_drop = nn.Dropout2d()

        self.conv5 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, padding=1)
        self.BN5 = nn.BatchNorm2d(num_features=64)
        self.conv6 = nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3, padding=1)
        self.BN6 = nn.BatchNorm2d(num_features=64)
        self.conv7 = nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3, padding=1)
        self.BN7 = nn.BatchNorm2d(num_features=64)
        self.conv8 = nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3, padding=1)
        self.BN8 = nn.BatchNorm2d(num_features=64)

        self.conv9 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, padding=1)
        self.BN9 = nn.BatchNorm2d(num_features=128)
        self.conv10 = nn.Conv2d(in_channels=128, out_channels=128, kernel_size=3, padding=1)
        self.BN10 = nn.BatchNorm2d(num_features=128)
        self.conv11 = nn.Conv2d(in_channels=128, out_channels=128, kernel_size=3, padding=1)
        self.BN11 = nn.BatchNorm2d(num_features=128)
        self.conv12 = nn.Conv2d(in_channels=128, out_channels=128, kernel_size=3, padding=1)
        self.BN12 = nn.BatchNorm2d(num_features=128)

        self.conv13 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3, padding=1)
        self.BN13 = nn.BatchNorm2d(num_features=256)
        self.conv14 = nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3, padding=1)
        self.BN14 = nn.BatchNorm2d(num_features=256)
        self.conv15 = nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3, padding=1)
        self.BN15 = nn.BatchNorm2d(num_features=256)
        self.conv16 = nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3, padding=1)
        self.BN16 = nn.BatchNorm2d(num_features=256)

        #self.fc1 = nn.Linear(16384, 8192)
        #self.fc2 = nn.Linear(8192, 4096)
        self.fc1 = nn.Linear(4096, 2048)
        self.fc2 = nn.Linear(2048, 1024)
        self.fc3 = nn.Linear(1024, 512)
        self.fc4 = nn.Linear(512, 200)

    def forward(self, x):
        x = self.conv2_drop(F.relu(self.BN1(self.conv1(x))))
        x = self.conv2_drop(F.relu(self.BN2(self.conv2(x))))
        x = self.conv2_drop(F.relu(self.BN3(self.conv3(x))))
        x = self.conv2_drop(F.relu(self.BN4(self.conv4(x))), 2))

        x = self.conv2_drop(F.relu(self.BN5(self.conv5(x))))
        x = self.conv2_drop(F.relu(self.BN6(self.conv6(x))))
        x = self.conv2_drop(F.relu(self.BN7(self.conv7(x))))
        x = self.conv2_drop(F.relu(self.BN8(self.conv8(x))), 2))

        x = self.conv2_drop(F.relu(self.BN9(self.conv9(x))))
        x = self.conv2_drop(F.relu(self.BN10(self.conv10(x))))
        x = self.conv2_drop(F.relu(self.BN11(self.conv11(x))))
        x = self.conv2_drop(F.relu(self.BN12(self.conv12(x))), 2))

        x = self.conv2_drop(F.relu(self.BN13(self.conv13(x))))
        x = self.conv2_drop(F.relu(self.BN14(self.conv14(x))))
        x = self.conv2_drop(F.relu(self.BN15(self.conv15(x))))
        x = self.conv2_drop(F.relu(self.BN16(self.conv16(x))), 2))

        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = F.relu(self.fc3(x))
        x = self.fc4(x)
        return x
```

Figure 7: Classifier 5 architecture

## Classifier 7 Architecture

Number 6 was skipped as there were two classifiers for number 4 (4 and 4.2). this classifier Adaptive Average pooling in between the convolutions that are based on Classifier 4 design and sums up these intermediary value when it returns the logit as  $x+x1+x2$ . This provides a much needed advantage of classifier 4.2 as it has a much smaller size after removing the FC layers that were previously put in 4.2.

```
class Classifier_deep_convs(nn.Module):
    def __init__(self):
        super(Classifier_deep_convs, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=128, kernel_size=3, padding=1)
        self.BN1 = nn.BatchNorm2d(num_features=128)
        self.conv2 = nn.Conv2d(in_channels=128, out_channels=128, kernel_size=3, padding=1)
        self.BN2 = nn.BatchNorm2d(num_features=128)

        self.conv2_drop = nn.Dropout2d()

        self.conv3 = nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3, padding=1)
        self.BN3 = nn.BatchNorm2d(num_features=256)
        self.conv4 = nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3, padding=1)
        self.BN4 = nn.BatchNorm2d(num_features=256)

        self.fcmid1 = nn.Linear(256, 200)
        self.conv5 = nn.Conv2d(in_channels=256, out_channels=512, kernel_size=3, padding=1)
        self.BN5 = nn.BatchNorm2d(num_features=512)
        self.conv6 = nn.Conv2d(in_channels=512, out_channels=512, kernel_size=3, padding=1)
        self.BN6 = nn.BatchNorm2d(num_features=512)

        self.fcmid2 = nn.Linear(512, 200)
        self.conv7 = nn.Conv2d(in_channels=512, out_channels=1024, kernel_size=3, padding=1)
        self.BN7 = nn.BatchNorm2d(num_features=1024)
        self.conv8 = nn.Conv2d(in_channels=1024, out_channels=1024, kernel_size=3, padding=1)
        self.BN8 = nn.BatchNorm2d(num_features=1024)

        self.fc1 = nn.Linear(1024, 512)
        self.fc2 = nn.Linear(512, 200)

    def forward(self, x):
        x = F.relu(self.BN1(self.conv1(x)))
        x = self.conv2_drop(F.max_pool2d(F.relu(self.conv2(x)), 2))

        x = self.conv2_drop(F.relu(self.BN3(self.conv3(x))))
        x = self.conv2_drop(F.relu(self.BN4(self.conv4(x))))

        x1 = F.adaptive_avg_pool2d(x, (1, 1))
        x1 = x1.view(x1.size(0), -1)
        x1 = self.fcmid1(x1) #1st clf

        x = F.relu(self.BN5(self.conv5(x)))
        x = self.conv2_drop(F.relu(self.BN6(self.conv6(x))))

        x2 = F.adaptive_avg_pool2d(x, (1, 1))
        x2 = x2.view(x2.size(0), -1)
        x2 = self.fcmid2(x2) #2nd clf

        x = F.relu(self.BN7(self.conv7(x)))
        x = self.conv2_drop(F.relu(self.BN8(self.conv8(x))))
        x = F.adaptive_avg_pool2d(x, (1, 1))
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        x = x+x1+x2
        return x
```

Figure 8: Classifier 7 architecture

## Data Augmentation Methods:

Three types of Data Augmentation methods were used. The first one was switching off some pixels (5%) in the images at random every epoch. Second method was to produce new image samples by rotating at 90, 45, 25 and 65 degrees angles and flipping horizontally. The following sections have more details.

## Switching off pixels

In this method it is like adding noise to the image. The 3-D tensor of 5% zeros was created at random every single epoch and then multiplied by the input images during the training. This method was not effective.


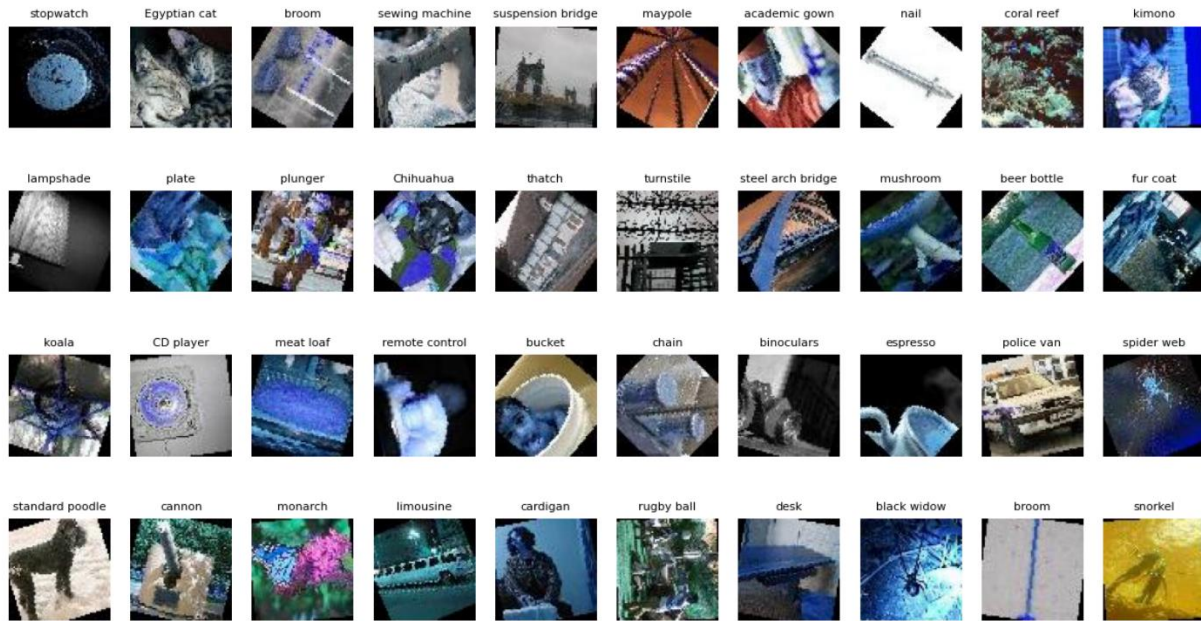


Figure 9: Data Augmentation, random switching off of pixel (noise)

## Creating new images

This method involved rotating and flipping the images. With the angles used and the flip the result was 4000 images per class. This sums up to 800k images for the training set. This made it extremely difficult to run for long epochs due to resources limitations (colab allows for limited hours). When tested with classifier 4.2 the accuracy was 10% at 4 epochs which is a good jump. However, the resources did not allow to train further.



*Figure 10: Data augmentation, rotation and flipping*

## Results

The results for training loss, test loss and accuracy were saved for all the models (except clf2 due to runtime restart). Below are the curves for the training loss for the classifier recorded every 100 iterations. Some classifiers were trained longer and others lesser than 30 epochs depending on the seen results during the training. It can be seen that resent classifier had the quickest drop in training loss.

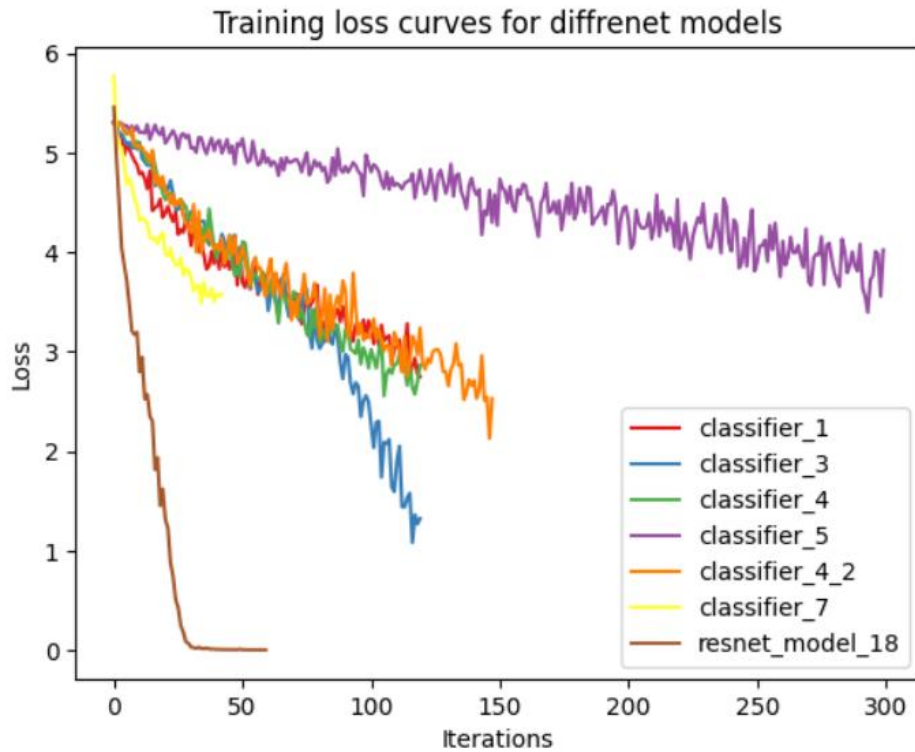


Figure 11: Training loss chart for all the models

The curves below show the test loss for the models. When the curve starts increasing like in the case of resnet, it can be confirmed that the model is over fitting.

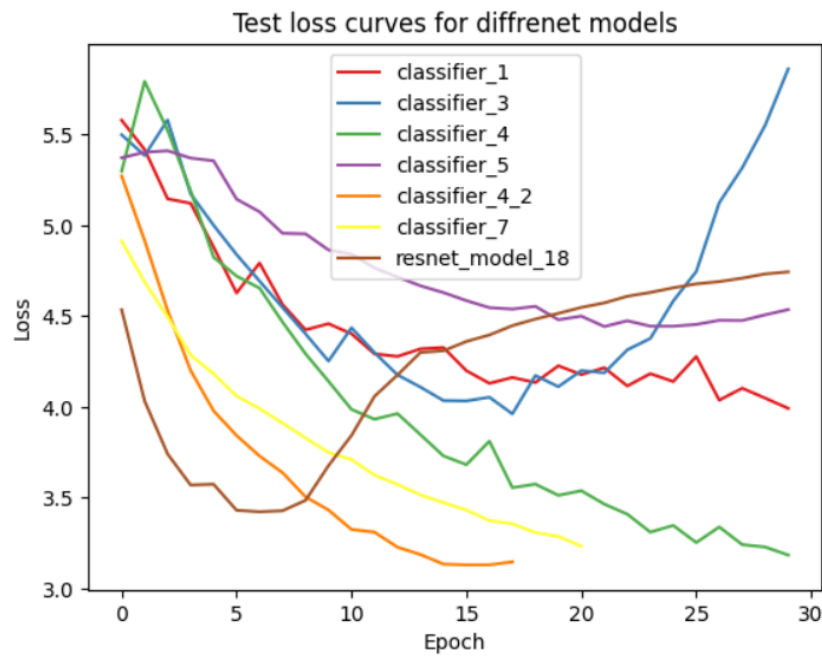


Figure 12: test loss for all models



The chart below shows the accuracy of the different models. Since model 4 showed some promising test accuracy results it was further developed with 4.2 and 7. These three classifiers (4, 4.2 and 7) performed better than resnet18 with delayed reach to results. However, from the test loss curves above they do not over fit like resnet18 did.




Figure 13: Models accuracy over epochs

Using a pretrained resnet18 model produce 53.5% accuracy over 20 epochs. This shows that the models designed in this project while they're less than 18 convolutions have achieved acceptable classification results over 200 classes.




Figure 14: resnet18 pretrained model accuracy




Figure 15: accuracy comparison for all the models

## Conclusions and recommendations

Classifying 200 classes can be a challenging task. To achieve high accuracy many models were constructed and tested. The models designed in this project are comparable with resnet18 model which shows acceptable results.

There were resources limitations to be able to train more models with different architectures which lead to not being able to use the data augmentation (800k training images) to improve the analysis.

Further work can be done with model design and the use of other neural networks (other than CNN) to analyze how different neural networks architectures can change improve the accuracy of classification.

## Project Hosting:

The project will be hosted in the following github repository

<https://github.com/m-hossni/INFO-6147-Deep-Learning.git>

## Final Notes

There are multiple notebooks in the github repo above. This is because different models were used at different stages in different notebooks in google colab.

You may find several sections of the code repeated like image display, this was done to ensure that the loader is assigned the correct labels.

You may also find some part of the training loop interrupted with error. This was done to avoid google colab erasing the variables and to be able to save the model parameters and logs.

Additionally, there are json logs for different models available in the github inside model\_data folder. The models parameter were also saved but due to their size they are not pushed to github.

## References:

Data source: [Tiny ImageNet](#)

[1] Data classes: [Tiny-Imagenet-200/sets/words200.txt at master · rmccorm4/Tiny-Imagenet-200](#)



## Appendix A: Tiny ImageNet Classes

n01443537	goldfish, Carassius auratus	n02099601	golden retriever
n01629819	European fire salamander, Salamandra	n02099712	Labrador retriever
salamandra		n02106662	German shepherd, German shepherd dog, German police dog, alsatian
n01641577	bullfrog, Rana catesbeiana	n02113799	standard poodle
n01644900	tailed frog, bell toad, ribbed toad, tailed toad, Ascaphus trui	n02123045	tabby, tabby cat
n01698640	American alligator, Alligator	n02123394	Persian cat
mississippiensis		n02124075	Egyptian cat
n01742172	boa constrictor, Constrictor constrictor	n02125311	cougar, puma, catamount, mountain lion, painter, panther, Felis concolor
n01768244	trilobite	n02129165	lion, king of beasts, Panthera leo
n01770393	scorpion	n02132136	brown bear, bruin, Ursus arctos
n01774384	black widow, Latrodectus mactans	n02165456	ladybug, ladybeetle, lady beetle, ladybird, ladybird beetle
n01774750	tarantula	n02190166	fly
n01784675	centipede	n02206856	bee
n01855672	goose	n02226429	grasshopper, hopper
n01882714	koala, koala bear, kangaroo bear, native bear, Phascolarctos cinereus	n02231487	walking stick, walkingstick, stick insect
n01910747	jellyfish	n02233338	cockroach, roach
n01917289	brain coral	n02236044	mantis, mantid
n01944390	snail	n02268443	dragonfly, darning needle, devil's darning needle, sewing needle, snake feeder, snake doctor, mosquito hawk, skeeter hawk
n01945685	slug	n02279972	monarch, monarch butterfly, milkweed butterfly, Danaus plexippus
n01950731	sea slug, nudibranch	n02281406	sulphur butterfly, sulfur butterfly
n01983481	American lobster, Northern lobster, Maine lobster, Homarus americanus	n02321529	sea cucumber, holothurian
n01984695	spiny lobster, langouste, rock lobster, crawfish, crayfish, sea crawfish	n02364673	guinea pig, Cavia cobaya
n02002724	black stork, Ciconia nigra	n02395406	hog, pig, grunter, squealer, Sus scrofa
n02056570	king penguin, Aptenodytes patagonica	n02403003	ox
n02058221	albatross, mollymawk	n02410509	bison
n02074367	dugong, Dugong dugon		
n02085620	Chihuahua		
n02094433	Yorkshire terrier		

n02415577	bighorn, bighorn sheep, cimarron, Rocky Mountain bighorn, Rocky Mountain sheep, Ovis canadensis	n02892201	brass, memorial tablet, plaque
n02423022	gazelle	n02906734	broom
n02437312	Arabian camel, dromedary, Camelus dromedarius	n02909870	bucket, pail
n02480495	orangutan, orang, orangutang, Pongo pygmaeus	n02917067	bullet train, bullet
n02481823	chimpanzee, chimp, Pan troglodytes	n02927161	butcher shop, meat market
n02486410	baboon	n02948072	candle, taper, wax light
n02504458	African elephant, Loxodonta africana	n02950826	cannon
n02509815	lesser panda, red panda, panda, bear cat, cat bear, Ailurus fulgens	n02963159	cardigan
n02666196	abacus	n02977058	cash machine, cash dispenser, automated teller machine, automatic teller machine, automated teller, automatic teller, ATM
n02669723	academic gown, academic robe, judge's robe	n02988304	CD player
n02699494	altar	n02999410	chain
n02730930	apron	n03014705	chest
n02769748	backpack, back pack, knapsack, packsack, rucksack, haversack	n03026506	Christmas stocking
n02788148	bannister, banister, balustrade, balusters, handrail	n03042490	cliff dwelling
n02791270	barbershop	n03085013	computer keyboard, keypad
n02793495	barn	n03089624	confectionery, confectionary, candy store
n02795169	barrel, cask	n03100240	convertible
n02802426	basketball	n03126707	crane
n02808440	bathtub, bathing tub, bath, tub	n03160309	dam, dike, dyke
n02814533	beach wagon, station wagon, wagon, estate car, beach waggon, station waggon, waggon	n03179701	desk
n02814860	beacon, lighthouse, beacon light, pharos	n03201208	dining table, board
n02815834	beaker	n03250847	drumstick
n02823428	beer bottle	n03255030	dumbbell
n02837789	bikini, two-piece	n03355925	flagpole, flagstaff
n02841315	binoculars, field glasses, opera glasses	n03388043	fountain
n02843684	birdhouse	n03393912	freight car
n02883205	bow tie, bow-tie, bowtie	n03400231	frying pan, frypan, skillet
		n03404251	fur coat
		n03424325	gasmask, respirator, gas helmet
		n03444034	go-kart
		n03447447	gondola

n03544143	hourglass	n04070727	refrigerator, icebox
n03584254	iPod	n04074963	remote control, remote
n03599486	jinrikisha, ricksha, rickshaw	n04099969	rocking chair, rocker
n03617480	kimono	n04118538	rugby ball
n03637318	lampshade, lamp shade	n04133789	sandal
n03649909	lawn mower, mower	n04146614	school bus
n03662601	lifeboat	n04149813	scoreboard
n03670208	limousine, limo	n04179913	sewing machine
n03706229	magnetic compass	n04251144	snorkel
n03733131	maypole	n04254777	sock
n03763968	military uniform	n04259630	sombrero
n03770439	miniskirt, mini	n04265275	space heater
n03796401	moving van	n04275548	spider web, spider's web
n03804744	nail	n04285008	sports car, sport car
n03814639	neck brace	n04311004	steel arch bridge
n03837869	obelisk	n04328186	stopwatch, stop watch
n03838899	oboe, hautboy, hautbois	n04356056	sunglasses, dark glasses, shades
n03854065	organ, pipe organ	n04366367	suspension bridge
n03891332	parking meter	n04371430	swimming trunks, bathing trunks
n03902125	pay-phone, pay-station	n04376876	syringe
n03930313	picket fence, paling	n04398044	teapot
n03937543	pill bottle	n04399382	teddy, teddy bear
n03970156	plunger, plumber's helper	n04417672	thatch, thatched roof
n03976657	pole	n04456115	torch
n03977966	police van, police wagon, paddy wagon, patrol wagon, wagon, black Maria	n04465501	tractor
n03980874	poncho	n04486054	triumphal arch
n03983396	pop bottle, soda bottle	n04487081	trolleybus, trolley coach, trackless trolley
n03992509	potter's wheel	n04501370	turnstile
n04008634	projectile, missile	n04507155	umbrella
n04023962	punching bag, punch bag, punching ball, punchball	n04532106	vestment
n04067472	reel	n04532670	viaduct
		n04540053	volleyball

n04560804	water jug
n04562935	water tower
n04596742	wok
n04597913	wooden spoon
n06596364	comic book
n07579787	plate
n07583066	guacamole
n07614500	ice cream, icecream
n07615774	ice lolly, lolly, lollipop, popsicle
n07695742	pretzel
n07711569	mashed potato
n07715103	cauliflower
n07720875	bell pepper
n07734744	mushroom
n07747607	orange
n07749582	lemon
n07753592	banana
n07768694	pomegranate
n07871810	meat loaf, meatloaf
n07873807	pizza, pizza pie
n07875152	potpie
n07920052	espresso
n09193705	alp
n09246464	cliff, drop, drop-off
n09256479	coral reef
n09332890	lakeside, lakeshore
n09428293	seashore, coast, seacoast, sea-coast
n12267677	acorn